



Per-topic profile cards — Salinas-A corrected, K = 6. Each panel: $\phi_k(\lambda)$ with top-5 word λ -marks (red dashed); title = dominant label + per-topic prevalence + dominant-doc count.


t0 · prev 8.3% · N=66
class_14 70%




t1 · prev 12.0% · N=120
class_10 50%




t2 · prev 21.3% · N=218
Broccoli_green_weed 100%



t3 · prev 22.1% · N=249
class_13 88%



t4 · prev 8.1% · N=118
class_1 49%



t5 · prev 27.4% · N=299
class_13 100%

